

Clusters and Heatmaps

Jeff Oliver

2026-02-12

Want to combine the visualization of quantitative data with clustering algorithms? Unsatisfied with the options provided by the base R packages? In this hands-on workshop, we'll use the **ggendro** package for R to make publication-quality graphics.

Learning objectives

1. Run clustering analysis and display dedrogram
 2. Visualize data in a heatmap
 3. Use **grid** package to create multi-plot figures
-

Getting started

First we need to setup our development environment. Open RStudio and create a new project via:

- File > New Project...
- Select 'New Directory'
- For the Project Type select 'New Project'
- For Directory name, call it something like "r-graphing" (without the quotes)
- For the subdirectory, select somewhere you will remember (like "My Documents" or "Desktop")

We need to create two folders: 'data' will store the data we will be analyzing, and 'output' will store the results of our analyses. In the RStudio console:

```
dir.create("data")  
dir.create("output")
```

We will also need a few R packages that are not included in the standard distribution of R. Those packages are `ggplot2`, `ggdendro`, `tidyr`, and `grid` and can be installed with the `install.packages()` function. Note these packages need only be installed once on your machine.

```
install.packages("ggplot2")
install.packages("ggdendro")
install.packages("grid")
install.packages("tidyr")
```

Finally, download data file from <https://jcoliver.github.io/learn-r/data/otter-mandible-data.csv> or <http://tinyurl.com/otter-data-csv> (the latter just re-directs to the former). These data are a subset of those used in a study on skull morphology and diet specialization in otters [doi: 10.1371/journal.pone.0143236](https://doi.org/10.1371/journal.pone.0143236). **Save this file in the “data” folder we just created.**

Data preparation

For any work that we do, we want to record all the steps we take, so instead of typing commands directly into the R console, we keep all our work in an R script. These scripts are just text files with R commands; by convention, we start the script with a brief bit of information about what the script does.

```
# Cluster & heatmap on otter data
# Jeff Oliver
# jcoliver@email.arizona.edu
# 2017-08-15

# Load dependencies
library("ggplot2")
library("ggdendro")
library("tidyr")
library("grid")

# Read in data
otter <- read.csv(file = "data/otter-mandible-data.csv",
                  stringsAsFactors = TRUE)
```

If we look at the first few rows of data, we see the first three columns have information about the specimen from which the measurements were taken, while columns 4-9 have data for six mandible measurements (`m1` - `m6`):

```
head(otter)
```

	species	museum	accession	m1	m2	m3	m4	m5	m6
1	A. cinerea	AMNH	101458	15.100	27.790	21.885	13.010	10.500	61.635
2	A. cinerea	AMNH	101461	12.740	26.750	20.265	13.255	8.340	59.370
3	A. cinerea	AMNH	101466	12.425	25.915	20.735	12.300	9.430	56.270
4	A. cinerea	AMNH	101635	13.400	28.030	22.075	10.580	10.455	58.080
5	A. cinerea	AMNH	101459	14.400	26.160	21.385	12.100	9.600	58.635
6	A. cinerea	AMNH	101462	14.525	29.020	22.305	11.905	11.070	60.655

For our purposes, we only want to look at the data for two species, so we subset the data, including only those rows that have either “A. cinerea” or “L. canadensis” in the **species** column. Note, this *does not* alter the data in the original file we read into R; it only alters the data object **otter** currently in R’s memory.

```
two_species <- c("A. cinerea", "L. canadensis")
otter <- otter[otter$species %in% two_species, ]
```

The last thing we need to do is scale the data variables (columns 4-9) so measurements have a mean value of 0 and a variance of 1.

```
otter_scaled <- otter
otter_scaled[, c(4:9)] <- scale(otter_scaled[, 4:9])
```


Clustering

To make our figure, we will build the two plots (the cluster diagram and the heatmap) separately, then use the **grid** framework to put them together. We start by making the dendrogram (or cluster).

```
# Run clustering
otter_matrix <- as.matrix(otter_scaled[, -c(1:3)])
rownames(otter_matrix) <- otter_scaled$accession
otter_dendro <- as.dendrogram(hclust(d = dist(x = otter_matrix)))

# Create dendro
dendro_plot <- gg dendrogram(data = otter_dendro, rotate = TRUE)
```


```
# Preview the plot
print(dendro_plot)
```



We can make those labels a bit smaller:

```
dendro_plot <- dendro_plot + theme(axis.text.y = element_text(size = 6))

# Preview the plot
print(dendro_plot)
```



That's it for the cluster image (we'll make a few more changes later on).

Heatmap

We need to start by transforming data to a “long” format, where each row only contains values for one measurement. For example, our data started in this format:

	species	museum	accession	m1	m2
1	A. cinerea	AMNH	101458	-0.1274976	-0.7008434
2	A. cinerea	AMNH	101461	-1.4031387	-1.0374187

We want to transform it so there is only a single measurement in each row:

```
# A tibble: 4 x 5
  species    museum accession measurement  value
  <fct>      <fct>   <fct>      <chr>         <dbl>
1 A. cinerea AMNH    101458      m1         -0.127
2 A. cinerea AMNH    101458      m2         -0.701
3 A. cinerea AMNH    101461      m1         -1.40
4 A. cinerea AMNH    101461      m2         -1.04
```

This new data has columns that have information about the specimen (the “species”, “museum”, and “accession” columns), but now has “measurement” and “value” columns. Instead of one column for each mandible measurement, each row only has one measurement in the “value” column, and the measurement type is stored in the “measurement” column.


We use the `tidyr` package’s `pivot_longer` function to handle the data transformation:

```
# Heatmap
# Data wrangling
otter_long <- pivot_longer(data = otter_scaled,
                           cols = -c(species, museum, accession),
                           names_to = "measurement",
                           values_to = "value")
```

Now we can use that new data frame, `otter_long` to create our heatmap. In the `ggplot` package, we use the `geom_tile` layer for creating a heatmap. Given our prior experience with the y-axis labels being large, we will again use `theme` to make the accession numbers (the y-axis labels) a little smaller:

```
heatmap_plot <- ggplot(data = otter_long, aes(x = measurement, y = accession)) +
  geom_tile(aes(fill = value)) +
  scale_fill_gradient2() +
  theme(axis.text.y = element_text(size = 6))


# Preview the heatmap
print(heatmap_plot)
```



Putting it all together

OK. We made a dendrogram. We made a heatmap. Now how to get them into the same figure? The `grid` package offers a way to combine separate plots.

```
grid.newpage()
print(heatmap_plot,
      vp = viewport(x = 0.4, y = 0.5, width = 0.8, height = 1.0))
print(dendro_plot,
      vp = viewport(x = 0.90, y = 0.445, width = 0.2, height = 1.0))
```



OK, that's a start, but there are a few things we need to address:

1. Notice the tips of the dendrogram are not in the same order as the y-axis of the heatmap. We'll need to re-order the heatmap to match the structure of the clustering.
2. It would be nice to have the dendrogram and heatmap closer together, without a legend separating them.
3. The dendrogram should be vertically stretched so each tip lines up with a row in the heatmap plot.

Re-order heatmap rows to match dendrogram

We'll use the order of the tips in the dendrogram to re-order the rows in our heatmap. The order of those tips are stored in the `otter_dendro` object, and we can extract it with the `order.dendro` function:

```
otter_order <- order.dendrogram(otter_dendro)
```

And now for the fun part. By default, `ggplot` use the level order of the y-axis labels as the means of ordering the rows in the heatmap. That is, level 1 of the factor is plotted in the top row, level 2 is plotted in the second row, level 3 in the third row and so on. So we can dictate the order in which these rows are plotted by *setting* the order of the levels in the column we use for the y-axis labels in the heatmap. The data we used for the heatmap was the long-format,


`otter_long` object, and the column we used for the y-axis is `accession`. Using the order of the dendrogram tips (we stored this above in the `otter_order` vector), we can re-level the `otter_long$accession` column to match the order of the tips in the trees.

```
# Order the levels according to their position in the cluster
otter_long$accession <- factor(x = otter_long$accession,
                              levels = otter_scaled$accession[otter_order],
                              ordered = TRUE)
```

At this point, we need to re-create the heatmap, because the underlying data (`otter_long`) has changed:

```
# Heatmap
heatmap_plot <- ggplot(data = otter_long, aes(x = measurement, y = accession)) +
  geom_tile(aes(fill = value)) +
  scale_fill_gradient2() +
  theme(axis.text.y = element_text(size = 6))


# Preview the heatmap
print(heatmap_plot)
```



Now the rows have been re-arranged and we can see right away that the heatmap now does a better job of grouping “red” rows together and “blue” rows together.

The combined figure shows the accession numbers in the heatmap are now in the same order as the accession numbers in the dendrogram:

```
grid.newpage()
print(heatmap_plot,
      vp = viewport(x = 0.4, y = 0.5, width = 0.8, height = 1.0))
print(dendro_plot,
      vp = viewport(x = 0.90, y = 0.445, width = 0.2, height = 1.0))
```




Re-position legend

We would like the legend to be somewhere else besides in between the two graphs, so we will move it to the top of the heatmap. We do this with through the `theme` layer when creating the heatmap, setting the value of `legend.position` to “top”:


```
# Heatmap
heatmap_plot <- ggplot(data = otter_long, aes(x = measurement, y = accession)) +
  geom_tile(aes(fill = value)) +
  scale_fill_gradient2() +
  theme(axis.text.y = element_text(size = 6),
        legend.position = "top")
```

```
# Preview the heatmap
print(heatmap_plot)
```



The combined figure now has the heatmap and dendrogram closer together:

```
grid.newpage()
print(heatmap_plot,
      vp = viewport(x = 0.4, y = 0.5, width = 0.8, height = 1.0))
print(dendro_plot,
      vp = viewport(x = 0.90, y = 0.445, width = 0.2, height = 1.0))
```



Align dendrogram tips with heatmap rows

Finally, we will need to adjust the height and alignment of dendrogram to ensure the tips line up with the corresponding rows in the heatmap. Unfortunately, because the two figures are actually independent graphics objects, this part involves considerable trial and error. In general, we will need to change two parameters in the `viewport` function when we call `print` on `dendro_plot`:


1. `y`: This argument sets the vertical justification. A value of 0.5 centers the graphic in the middle of the viewport to which it is assigned. Values closer to zero move the graphic towards the bottom of the viewport and values closer to one move the graphic towards the top of the viewport. In our case, we will probably want to move the dendrogram down a little bit, so we use a value below 0.5.
2. `height`: This is the proportion of the viewport's vertical space the graphic should use. We initially set it to use all the vertical space (`height = 1.0`), but since the legend now occupies the real estate at the top of the graphic, we will need to make the dendrogram height smaller.

Note: The actual values you use will depend on the graphics device you ultimately use; that is, the values presented here are designed for an html to be viewed on the web, but if you are displaying on the screen or writing to a pdf file, you may need to use different values for `y` and `height` (e.g. when calling `print` on `dendro_plot`, it may work better on your screen to set `y = 0.4` and `height = 0.85`).

```

grid.newpage()
print(heatmap_plot,
      vp = viewport(x = 0.4, y = 0.5, width = 0.8, height = 1.0))
print(dendro_plot,
      vp = viewport(x = 0.90, y = 0.43, width = 0.2, height = 0.92))

```



And lastly, as a bonus, we can remove the accession numbers from the heatmap, as they are already displayed at the tips of the dendrogram (and we made sure the heatmap rows and dendrogram tips match, above). We do so by setting the relevant theme elements to `element_blank()` in our call to `ggplot`:


```

# Heatmap
heatmap_plot <- ggplot(data = otter_long, aes(x = measurement, y = accession)) +
  geom_tile(aes(fill = value)) +
  scale_fill_gradient2() +
  theme(axis.text.y = element_blank(),
        axis.title.y = element_blank(),
        axis.ticks.y = element_blank(),
        legend.position = "top")

grid.newpage()
print(heatmap_plot,
      vp = viewport(x = 0.4, y = 0.5, width = 0.8, height = 1.0))

```

```
print(dendro_plot,
      vp = viewport(x = 0.90, y = 0.43, width = 0.2, height = 0.92))
```



Our final script would look like this:

```
# Cluster & heatmap on otter data
# Jeff Oliver
# jcoliver@email.arizona.edu
# 2017-08-15

#####

# Load dependencies
library("ggplot2")
library("ggdendro")
library("tidyr")
library("grid")

# Read in data
otter <- read.csv(file = "data/otter-mandible-data.csv",
                  stringsAsFactors = TRUE)

# Restrict the data to only two species, A. cinerea & L. canadensis
```

```

two_species <- c("A. cinerea", "L. canadensis")
otter <- otter[otter$species %in% two_species, ]

# Force re-numbering of the rows after subsetting
rownames(otter) <- NULL

# Scale each measurement (independently) to have a mean of 0 and variance of 1
otter_scaled <- otter
otter_scaled[, c(4:9)] <- scale(otter_scaled[, 4:9])

# Run clustering
otter_matrix <- as.matrix(otter_scaled[, -c(1:3)])
rownames(otter_matrix) <- otter_scaled$accession
otter_dendro <- as.dendrogram(hclust(d = dist(x = otter_matrix)))

# Create dendrogram plot
dendro_plot <- gg dendrogram(data = otter_dendro, rotate = TRUE) +
  theme(axis.text.y = element_text(size = 6))

# Heatmap

# Data wrangling
otter_long <- pivot_longer(data = otter_scaled,
                           cols = -c(species, museum, accession),
                           names_to = "measurement",
                           values_to = "value")

# Extract the order of the tips in the dendrogram
otter_order <- order.dendrogram(otter_dendro)
# Order the levels according to their position in the cluster
otter_long$accession <- factor(x = otter_long$accession,
                              levels = otter_scaled$accession[otter_order],
                              ordered = TRUE)

# Create heatmap plot
heatmap_plot <- ggplot(data = otter_long, aes(x = measurement, y = accession)) +
  geom_tile(aes(fill = value)) +
  scale_fill_gradient2() +
  theme(axis.text.y = element_blank(),
        axis.title.y = element_blank(),
        axis.ticks.y = element_blank(),
        legend.position = "top")

```

```
# All together
grid.newpage()
print(heatmap_plot,
      vp = viewport(x = 0.4, y = 0.5, width = 0.8, height = 1.0))
print(dendro_plot,
      vp = viewport(x = 0.90, y = 0.43, width = 0.2, height = 0.92))
```

Additional resources

- Documentation for the [ggdendro package](#)
- An [example](#) on StackOverflow
- A [pdf version](#) of this lesson